

Data Analyst Assignment

Rebecca Cockroft

2023-01-16

Problem 1 (Merge)

Combining New York and Detroit purchases:

```
# loading required packages

library(tidyverse)

## --- Attaching packages --- tidyverse 1.3.2 ---
## ✓ ggplot2 3.4.0      ✓ purrr   1.0.1
## ✓ tibble  3.1.8      ✓ dplyr   1.0.10
## ✓ tidyr   1.2.1      ✓ stringr 1.5.0
## ✓ readr   2.1.3      ✓ forcats 0.5.2
## --- Conflicts --- tidyverse_conflicts() ---
## * dplyr::filter() masks stats::filter()
## * dplyr::lag()     masks stats::lag()

# importing csv files

ny <- read_csv("new_york_purchases.csv")

## Rows: 27 Columns: 5
## --- Column specification ---
## Delimiter: ",",
## chr (3): barcode, purchase_timestamp, type
## dbl (2): id, amount
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

detroit <- read_csv("detroit_purchases.csv")

## Rows: 27 Columns: 5
## --- Column specification ---
## Delimiter: ",",
## chr (3): barcode, amount, type
## dbl (1): id
## dtm (1): purchase_timestamp
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

# not required but it might be a good idea to create another "location" variable
# by doing so we can tell which dataset it originally came from, and allows for future filtering

ny$location <- "New York"
detroit$location <- "Detroit"

# let's check the data types to make sure they are the same for ny and detroit

str(detroit)

## spc_tbl_ [27 × 6] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ id          : num [1:27] 0 1 2 3 4 5 6 7 8 9 ...
## $ barcode     : chr [1:27] "1835566385273" "7758948061357" "7410144862762" "3470282535786" ...
## $ amount      : chr [1:27] "$1.61" "$3.86" "$2.56" "$2.3" ...
## $ purchase_timestamp: POSIXct[1:27], format: "2023-01-01 02:58:07" "2023-01-01 20:34:02" ...
## $ type        : chr [1:27] "vegetable" "vegetable" "dairy" "vegetable" ...
## $ location     : chr [1:27] "Detroit" "Detroit" "Detroit" "Detroit" ...
## - attr(*, "spec")=
## .. cols(
## .. id = col_double(),
## .. barcode = col_character(),
## .. amount = col_character(),
## .. purchase_timestamp = col_datetime(format = ""),
## .. type = col_character()
## .. )
## - attr(*, "problems")=<externalptr>

str(ny)

## spc_tbl_ [27 × 6] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ id          : num [1:27] 0 1 2 3 4 5 6 7 8 9 ...
## $ barcode     : chr [1:27] "0766635859597" "1170285435077" "6963387314679" "5357547177449" ...
## $ amount      : num [1:27] 3.01 3.48 3.66 3.07 3.74 2.86 2.25 3.81 3.29 2.43 ...
## $ purchase_timestamp: chr [1:27] "2023-01-01 08:33:37 +0000" "2023-01-01 00:41:34 +0000" "2023-01-01 18:22:27 +0000" "2023-01-01 12:55:49 +0000" ...
## $ type        : chr [1:27] "puffs" "cakes" "tomato" "beans" ...
## $ location     : chr [1:27] "New York" "New York" "New York" "New York" ...
## - attr(*, "spec")=
## .. cols(
## .. id = col_double(),
## .. barcode = col_character(),
## .. amount = col_double(),
## .. purchase_timestamp = col_character(),
## .. type = col_character()
## .. )
## - attr(*, "problems")=<externalptr>

# NY timestamp is not in date/time (POSIX) format, let's change that
# Detroit purchases are automatically coded as UTC when read in, let's change this to EST to be consistent with N Y.

detroit$purchase_timestamp <- as.POSIXct(as.character(detroit$purchase_timestamp), tz = "EST")
ny$purchase_timestamp <- as.POSIXct(ny$purchase_timestamp, tz = "EST")

# make amount numeric for detroit

detroit$amount <- parse_number(detroit$amount)

# merging the datasets by appending them to each other

full_data <-
  rbind(detroit, ny) %>%

  # dropping "id" column because it's redundant to R's index and is misaligned after merge

  select(-id)

# previewing data

head(full_data)
```

## # A tibble: 6 × 5					
##	barcode	amount	purchase_timestamp	type	location
##	<chr>	<dbl>	<dtm>	<chr>	<chr>
## 1	1835566385273	1.61	2023-01-01 02:58:07	vegetable	Detroit
## 2	7758948061357	3.86	2023-01-01 20:34:02	vegetable	Detroit
## 3	7410144862762	2.56	2023-01-01 14:01:41	dairy	Detroit
## 4	3470282535786	2.3	2023-01-01 01:50:27	vegetable	Detroit
## 5	5583888078408	3.67	2023-01-01 21:49:34	dairy	Detroit
## 6	6986147139275	2.8	2023-01-01 17:23:27	dairy	Detroit

Normalizing type field to a product line:

```
# selecting rows with specific values and replacing "type" according to mapping

full_data$type[full_data$type %in% c("cakes", "pizzas", "puffs")] <- "bakery"
full_data$type[full_data$type %in% c("milk", "cheese")] <- "dairy"
full_data$type[full_data$type %in% c("tomato", "carrot", "beans")] <- "vegetable"

# checking that only bakery, dairy, and vegatable remain in the data

table(full_data$type)
```

##	bakery	dairy	vegetable
##	21	15	18

Problem 2 (Filter)

```
# only selecting rows with timestamp on 1/2/2023

jan_two <- full_data %>%
  filter(as.Date(as.character(purchase_timestamp)) == as.Date("2023-01-02"))
```

Problem 3 (Graph)

i. Total Revenue by Product Line

```
# creating barplot

ggplot()

# aggregating data to sum revenue by product line

jan_two %>%
  group_by(type) %>%
  summarize(amount = sum(amount)),

# setting x-axis and color to product line, y to revenue

aes(x=type, y=amount, fill = type)) +
  geom_bar(stat="identity") +

# relabeling categories and setting color palette

scale_fill_brewer(palette = "Set3",
  labels = c('Bakery', 'Dairy', 'Vegetable')) +
  scale_x_discrete(labels = c('Bakery', 'Dairy', 'Vegetable')) +

# putting dollar labels on y-axis

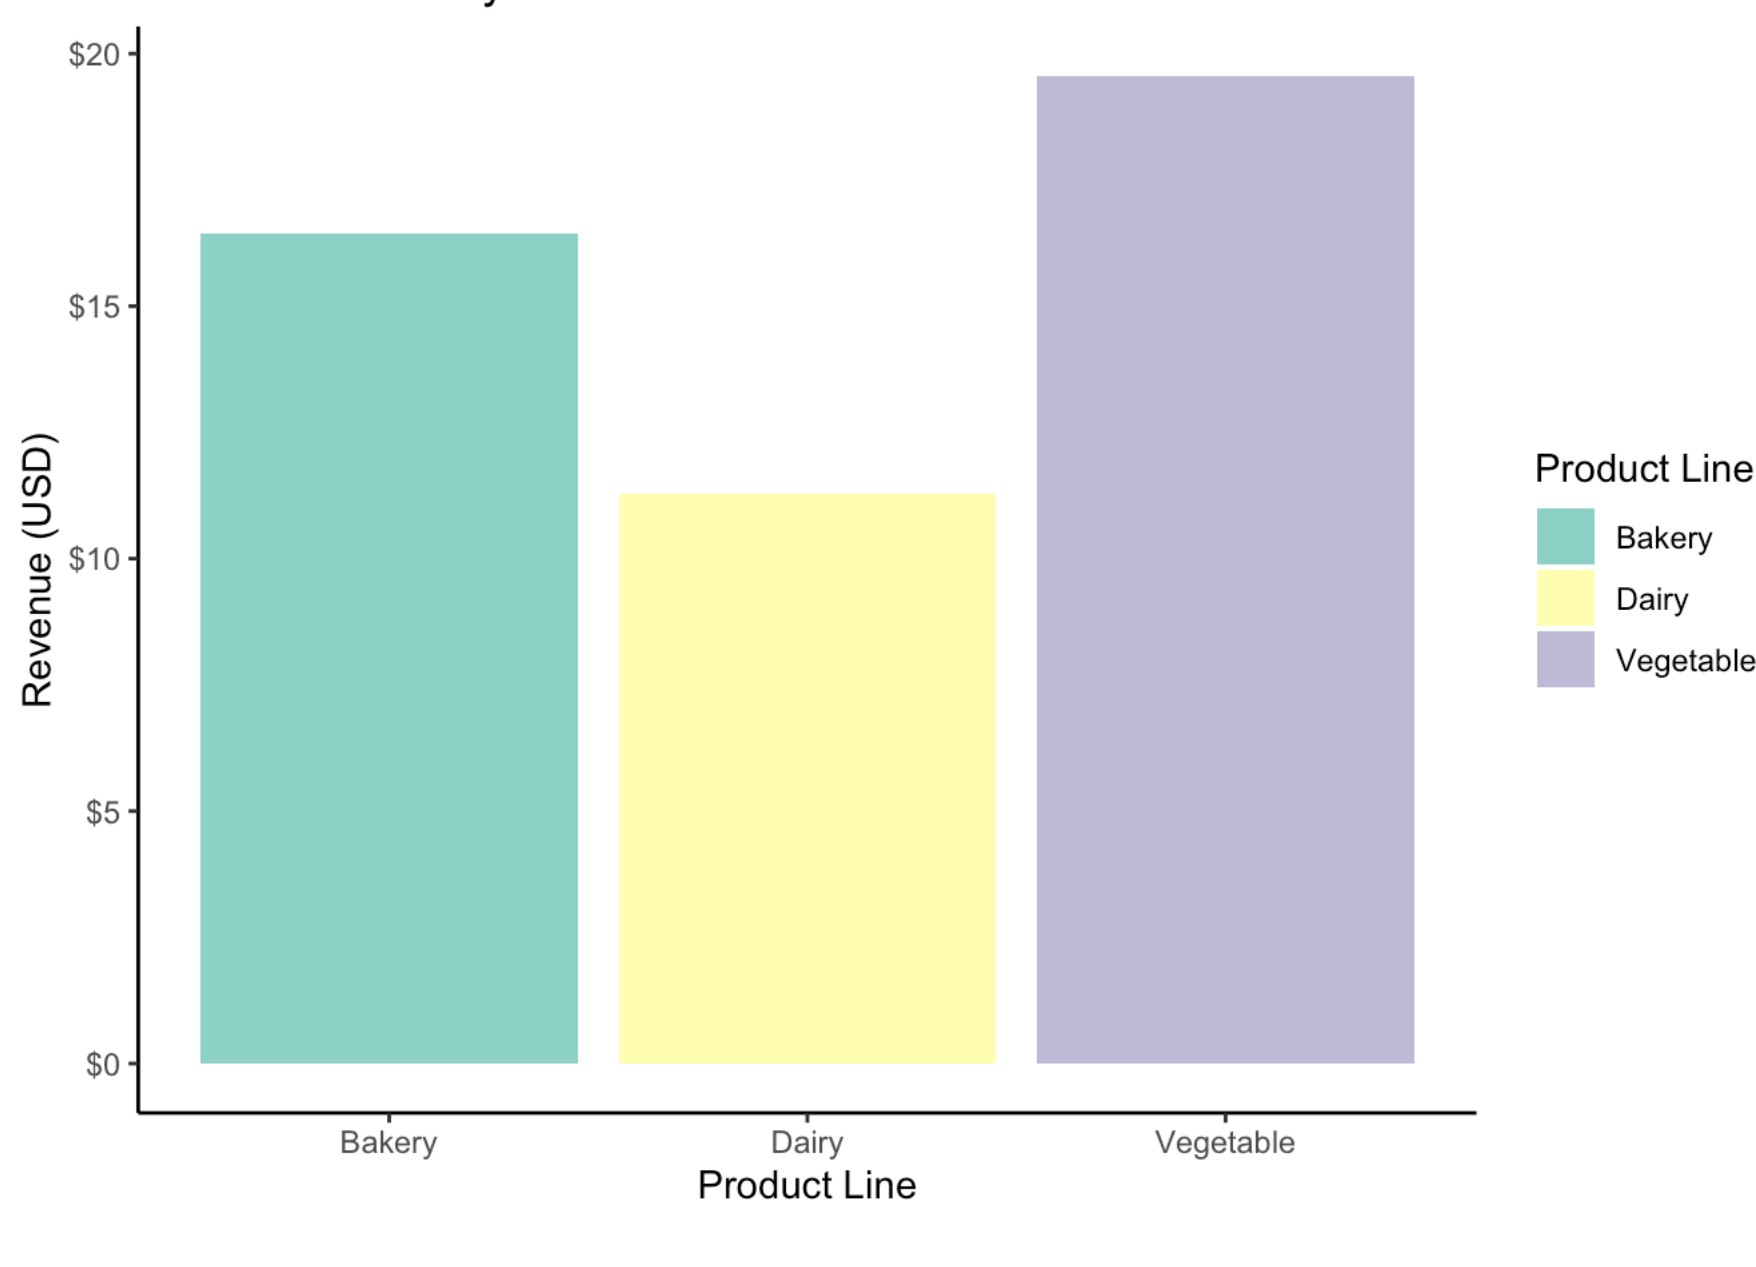
scale_y_continuous(labels=scales::dollar_format()) +

# setting legend/x-axis/y-axis/title labels

labs(fill = "Product Line",
  x = "Product Line",
  y = "Revenue (USD)",
  title = "Total Revenue by Product Line on 1/2/2023") +

# setting theme

theme_classic()
```



ii. Total Number of Items Purchased by Hour

```
# creating histogram

ggplot(jan_two,

# setting x-axis to purchase time

aes(x = purchase_timestamp)) +

# making the bin size 1 hour and setting colors

geom_histogram(binwidth = 3600,
  boundary = 0,
  fill = "lightblue",
  color = "black") +

# creating labels for each hour

scale_x_datetime(date_breaks = "1 hour",
  date_labels = "%H:%M",
  expand = c(.05, 0, .01, 0)) +

# setting x-axis/y-axis/title labels

labs(x = "Time",
  y = "Number of Purchases",
  title = "Purchases by Hour on 1/2/2023") +

# setting theme

theme_bw() +

# rotating x-axis labels so they don't overlap

theme(axis.text.x = element_text(angle = 60, hjust=1))
```

